

30.8 A 16nm, 1Mb, 1-to-8b-Configurable 444.21TOPS/W Fully Digital SRAM Compute-In-Memory Macro for Hybrid SNN-CNN Edge Computing

Yao-Kai Yeh^{*1}, Jian-Wei Su^{*2}, Ting-Hao Hsu¹, Mai Tseng¹, Jen-Chun Tien¹, Ko-Chi Chen¹, Chih-Yen Yue¹, Yu-En Lin¹, Yu-Jia Hu¹, Le-Jung Hsieh¹, Jyun-Cheng Bai¹, Yu-Sheng Kao¹, Tsung-Han Lou¹, Hung-Hsi Hsu¹, De-Qi You¹, Sheu Shyh-Shyuan², Wei-Chung Lo², Shih-Chieh Chang^{1,2}, Ya-Tang Yang¹, Chung-Chuan Lo¹, Ren-Shuo Liu¹, Chih-Cheng Hsieh¹, Kea-Tiong Tang¹, Meng-Fan Chang¹

¹National Tsing Hua University, Hsinchu, Taiwan, ²Industrial Technology Research Institute, Hsinchu, Taiwan

^{*}Equally Credited Authors (ECAs)

Abstract

We present a silicon-verified fully digital SRAM-CIM macro that supports multi-bit hybrid SNN-CNN processing. Key innovations include: (1) a compact CFD-SC-RC architecture that shares SNN-CNN MAC hardware to minimize area and energy overhead while enabling adaptive model and precision selection; and (2) a PS-WR-ODM scheme that balances partial-

sum storage and weight-update energy. Fabricated in 16nm FinFET CMOS, the test chip achieves 444.21TOPS/W (1bIN-8bW-14bOUT, SNN) and 62.84TOPS/W (8bIN-8bW-22bOUT, CNN).

Spiking neural networks (SNNs) [1] have demonstrated superior energy efficiency in event-driven applications such as event-based vision and audio processing, while convolutional neural networks (CNNs) [2-3] provide high accuracy in image and voice recognition tasks at the edge. Hybrid SNN-CNN models [4-6] have shown promise in balancing accuracy and compute energy, enabling robust edge inference for various applications. Meanwhile, SRAM-based digital compute-in-memory (SRAM-DCIM) [7] has been widely explored to reduce data movement and improve energy efficiency, while outperforming analog SRAM-CIM (SRAM-ACIM) [8] in MAC accuracy. Combining SRAM-DCIM with hybrid SNN-CNN models is a promising solution for diverse edge-AI tasks.

However, as shown in Fig. 30.8.1, designing SRAM-DCIM macros supporting hybrid SNN-CNN introduces two key challenges: (1) extra hardware overhead for supporting both SNN and CNN, and excessive compute energy in the digital domain; and (2) area-energy trade-off between partial sum (PSUM) storage and weight (W) update during input (IN) and W data mapping.

This work addresses these challenges with two key techniques: (1) a compact fully-digital SNN-CNN reconfigurable computing (CFD-SC-RC) architecture that shares SNN-CNN MAC circuit to reduce area and energy overhead, and supports adaptive model and precision selection across layers; and (2) a PSUM storage-weight reuse optimized data mapping (PS-WR-ODM) scheme that balances PSUM-storage minimization and weight-reuse maximization, achieving optimized register usage and weight-update energy. The test chip, fabricated in a foundry 16nm FinFET CMOS process, demonstrates 444.21 TOPS/W with a 1bIN-8bW-14bOUT configuration in SNN mode, and 62.84TOPS/W with an 8bIN-8bW-22bOUT configuration in CNN mode.

Figure 30.8.2 illustrates the proposed 1Mb CFD-SC-RC SRAM-DCIM macro, consisting of 16 64kb banks. Each bank integrates a control block, I/O block, IN/W scheduler, WL driver, 32 reconfigurable inference blocks (RIBs), an adder tree, and a digital multi-membrane potential dynamics behavior (MPDB) unit (DM-MPDBU). Each RIB contains a 2kb SRAM array, a weight stationary cell (WSC), eight cross-model dual-mode local computing cells (CMDM-LCCs), and a shift-and-add (SaA) circuit. Naïve SRAM-ACIM based hybrid SNN-CNN macros are vulnerable to process variations in MPDB functions, while prior SRAM-DCIM based hybrid SNN-CNN macros [4] support only a single MPDB type, limiting flexibility for diverse model requirements. Both architectures lack hardware sharing, incurring large area overhead. The proposed CFD-SC-RC architecture shares SNN-CNN MAC circuits to reduce area overhead, provides high tolerance to process variation, and, via the DM-MPDBU, additionally supports multiple MPDB functions beyond prior works [4], including integrate-and-fire (IF), leaky integrate-and-fire (LIF), and integer quadratic integrate-and-fire (IQIF) [9]. The IF mode provides simple spike generation; the LIF mode incorporates leakage dynamics for temporal noise filtering; and the IQIF mode introduces nonlinear integration for improved accuracy, enabling general-purpose SNN operations at the edge.

Figure 30.8.3 illustrates the CFD-SC-RC computation flow using the proposed CMDM-LCC and DM-MPDBU. Conventional bit-wise (BW)-IN multipliers (BWMs) for SNN cannot support word-wise (WW)-IN multiplication for CNN. To enable both BW-IN MAC for SNN and configurable WW-IN MAC for CNN, the proposed CMDM-LCC integrates two MUX-based multipliers (MULs), an INT8 toggle reducer (TR), dual-mode weight mapping (INT4/INT8) in the array, and an SaA circuit. Within one bank, each of eight CMDM-LCCs performs MAC operations with INs and Ws (W_{0-3}). In INT4 mode, the TR is disabled and four 4b MACs are executed per cycle: W_{0-1} map to two 4b weights ($A[3:0]$, $B[3:0]$) and W_{2-3} to another two ($C[3:0]$, $D[3:0]$). In INT8 mode, the TR eliminates redundant toggles, enabling two 8b MACs per cycle: W_{0-1} correspond to the four LSBs of $A[7:0]$ and $B[7:0]$, while W_{2-3} correspond to their four MSBs. The partial MACs ($pMAC_{0-7}$) are sent directly to the adder tree in SNN mode, while in CNN mode $pMAC_{0-7}$ are passed to the SaA circuit for WW-accumulation before entering the adder tree. Meanwhile, conventional MPDB blocks only support a single MPDB function, limiting flexibility for diverse SNN models. The proposed DM-MPDBU supports three functions: (1) IF mode—membrane potential (MP) accumulates until reaching threshold (peak), then fires and resets; (2) LIF mode—MP accumulates with leakage subtraction before proceeding as IF; and (3) IQIF mode—MP

accumulates with an integer quadratic (IQ) term dynamically defined by parameters (pdeth, rest, thresh, α , β): if $MP[t] < pdeth$, $IQ = (\alpha \times rest - MP[t]) >> 3$; otherwise, $IQ = (\beta \times (MP[t] - thresh)) >> 3$.

Figure 30.8.4 illustrates the proposed PS-WR-ODM scheme. In conventional timestep (TS)-first data mapping, all input neurons (IN_{0-k}) are accessed sequentially within each TS, minimizing PSUM storage but incurring high weight-update energy and long latency because weight reuse cannot be leveraged for early MPDB operations. In contrast, prior weight-reuse data mapping [10] accesses same-index INs (e.g., IN_0) across all TSs (TS_{0-N}), maximizing weight reuse and throughput but requiring large PSUM storage. The proposed PS-WR-ODM adopts a hybrid strategy in which TSs are grouped into blocks with user-defined size (e.g., pairs of adjacent TSs), and weight-reuse mapping is applied within each block. This approach achieves balanced performance in weight-update energy and latency with only marginal PSUM buffer overhead. For example, as shown in the timing diagrams, the conventional TS-first data mapping incurs long weight-load cycles, while the previous weight-reuse mapping demands excessive PSUM buffering. The proposed PS-WR-ODM achieves comparable latency reduction and energy savings without significant PSUM overhead.

Figure 30.8.5 shows the simulated performance of the proposed schemes. The CFD-SC-RC reduces macro area by 33% and improves energy efficiency by 2.31 $\times$ over prior hybrid SNN-CNN CIM design [4]. The PS-WR-ODM achieves >1.22 $\times$ improvement in FoM1, defined as: energy efficiency/(area $\times$ latency), compared to conventional timestep-first and weight-reuse-first [10] mappings. The combined implementation of CFD-SC-RC, CMDM-LCC, and PS-WR-ODM further delivers an FoM2, defined as inference accuracy $\times$ energy efficiency/(area $\times$ latency)), of 5.48 $\times$ over baseline designs [4,10].

Figure 30.8.6 presents the measurement results of the 16nm SRAM-DCIM macro and the position chart versus prior silicon-verified works. The Shmoo plot confirms 1.3ns computing latency (T_{MAC}) at 0.8V with an 8bIN-8bW-22bOUT configuration. The proposed macro achieves up to 8.2 $\times$ improvement in FoM3, defined as: IN precision $\times$ W precision $\times$ energy efficiency, over prior works. Figure 30.8.7 presents the die micrograph and chip performance summary. To the best of our knowledge, this is the first silicon-verified fully digital SRAM-CIM macro that supports multi-bit hybrid SNN-CNN processing.

Acknowledgement:

The authors would like to thank NTHU-ITRI collaboration and NSTC for their financial and manufacturing support.

References:

- [1] D. Huh and T. J. Sejnowski, "Gradient Descent for Spiking Neural Networks," *NeurIPS*, vol. 31, pp. 1–11, Dec. 2018. Accessed on Dec. 9, 2025. <https://proceedings.neurips.cc/paper_files/paper/2018/file/185e65bc40581880c4f2c82958de8cfe-Paper.pdf>
- [2] K. He et al., "Deep Residual Learning for Image Recognition," *IEEE Conf. Comp. Vision and Pattern Recognition (CVPR)*, pp. 770–778, Dec. 2016. <https://doi.org/10.1109/CVPR.2016.90>
- [3] M. Sandler et al., "MobileNetV2: Inverted Residuals and Linear Bottlenecks," *IEEE/CVF CVPR*, pp. 4510–4520, Dec. 2018. <https://doi.org/10.1109/CVPR.2018.00474>
- [4] C. Zha et al., "A Reconfigurable Digital Compute-In-Memory Heterogeneous Macro for Differential Frame Convolution and Spiking Neural Network," *ISCAS*, pp. 1–5, May 2025. <https://doi.org/10.1109/ISCAS56072.2025.11043929>
- [5] S. Kim et al., "C-DNN: A 24.5-85.8TOPS/W Complementary-Deep-Neural-Network Processor with Heterogeneous CNN/SNN Core Architecture and Forward-Gradient-Based Sparsity Generation," *ISSCC*, pp. 334–335, Feb. 2023. <https://doi.org/10.1109/ISSCC42615.2023.10067497>
- [6] S. Kim, S. Park, B. Na, and S. Yoon, "Spiking-YOLO: Spiking Neural Network for Energy-efficient Object Detection," *AAAI Conference on Artificial Intelligence*, vol. 34, no. 7, pp. 11270–11277, Feb. 2020. <https://doi.org/10.1609/aaai.v34i07.6787>

- [7] H. Fujiwara et al., "A 3nm, 32.5TOPS/W, 55.0TOPS/mm² and 3.78Mb/mm² Fully-Digital Compute-in-Memory Macro Supporting INT12 $\times$ INT12 with a Parallel-MAC Architecture and Foundry 6T-SRAM Bit Cell," *ISSCC*, pp. 572-573, Feb. 2024. <https://doi.org/10.1109/ISSCC49657.2024.10454556>
- [8] P. Chen et al., "A 22nm Delta-Sigma Computing-In-Memory ($\Delta\Sigma$ CIM) SRAM Macro with Near-Zero-Mean Outputs and LSB-First ADCs Achieving 21.38TOPS/W for 8b-MAC Edge AI Processing," *ISSCC*, pp. 140-141, Feb. 2023. <https://doi.org/10.1109/ISSCC42615.2023.10067289>
- [9] W.-C. Wu et al., "Integer Quadratic Integrate-and-Fire (IQIF): A Neuron Model for Digital Neuromorphic Systems," *AICAS*, pp. 1-4, June 2021. <https://doi.org/10.1109/AICAS51828.2021.9458572>
- [10] S. Kim et al., "SNPU: An Energy-Efficient Spike Domain Deep-Neural-Network Processor With Two-Step Spike Encoding and Shift-and-Accumulation Unit," *JSSC*, vol. 58, no. 10, pp. 2812-2825, Oct. 2023. <https://doi.org/10.1109/JSSC.2023.3270442>
- [11] K. Park et al., "A 701.7 TOPS/W Compute-in-Memory Processor With Time-Domain Computing for Spiking Neural Network," *IEEE Transactions on Circuits and Systems I: Regular Papers*, vol. 72, no. 1, pp. 25-35, Jan. 2025. <https://doi.org/10.1109/TCSI.2024.3480350>
- [12] H. Fu et al., "NeuC-CIM: A 1.3pJ/SOP Neuromorphic Charge-Domain Compute-in-Memory Macro for Spiking Neural Network," *IEEE Symp. VLSI Circuits*, pp. 1-3, June 2025. <https://doi.org/10.23919/VLSITechnologyandCirc6189.2025.11074927>




Figure 30.8.1: Motivation and challenges for hybrid SNN-CNN SRAM-CIM.




Figure 30.8.2: Structure of proposed SRAM-DCIM macro and CFD-SC-RC operation.




Figure 30.8.3: CDM-LCC and DM-MPDBU structure and operation.




Figure 30.8.4: PS-WR-ODM operational concept.




Figure 30.8.5: Simulated performance of proposed schemes.




Figure 30.8.6: Measurement results and position chart.




Figure 30.8.7: Die micrograph and performance summary table.